

Remote Procedure Call

- The need of communication
- Message transport and format
- Synchronous and asynchronous communication patterns
- Traditional load balancers
- Service Registry and Discovery
- Service meshes
- Idempotent service design

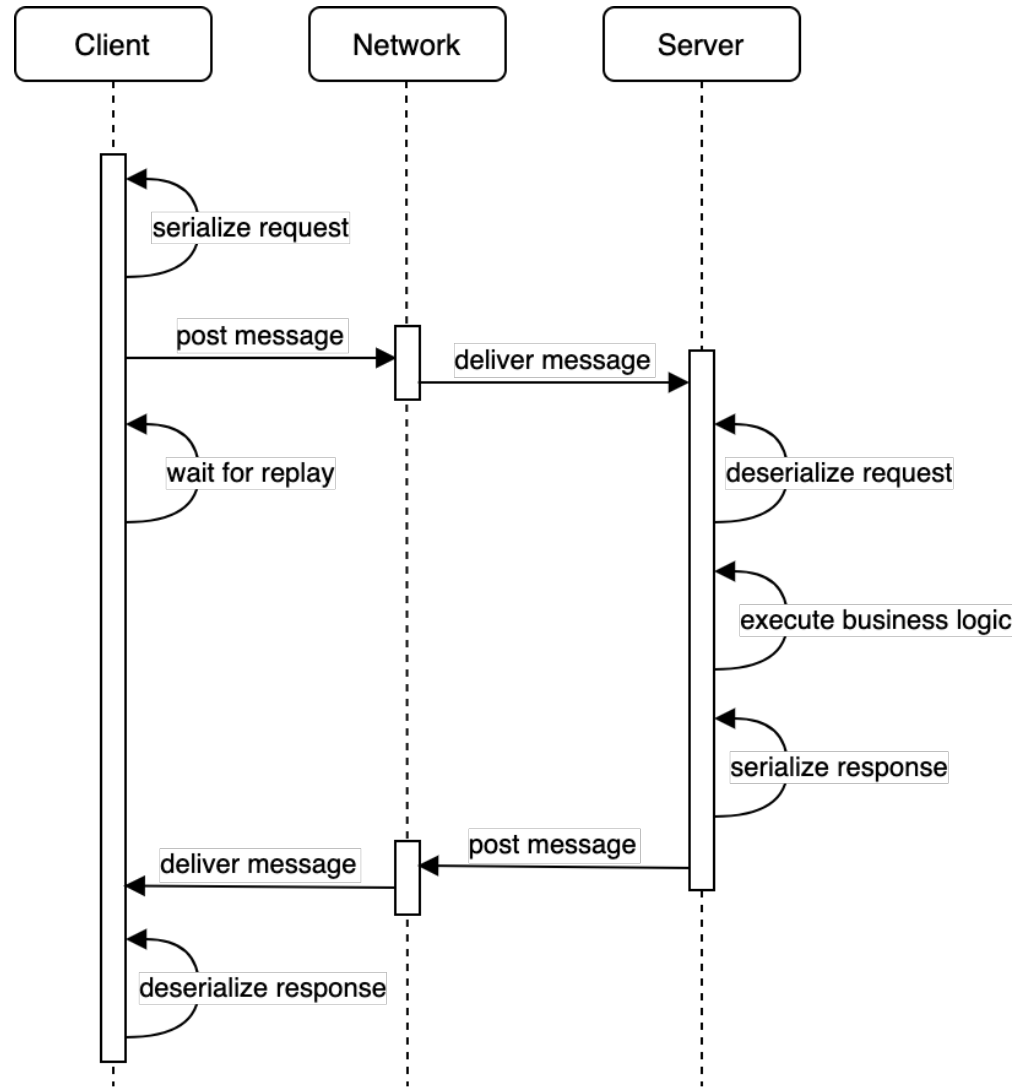
Definition

Distributed system is a computing environment, where multiple processes running on different machines, **communicate through the network** and coordinate actions in order to appear to the end-user as a single coherent system.



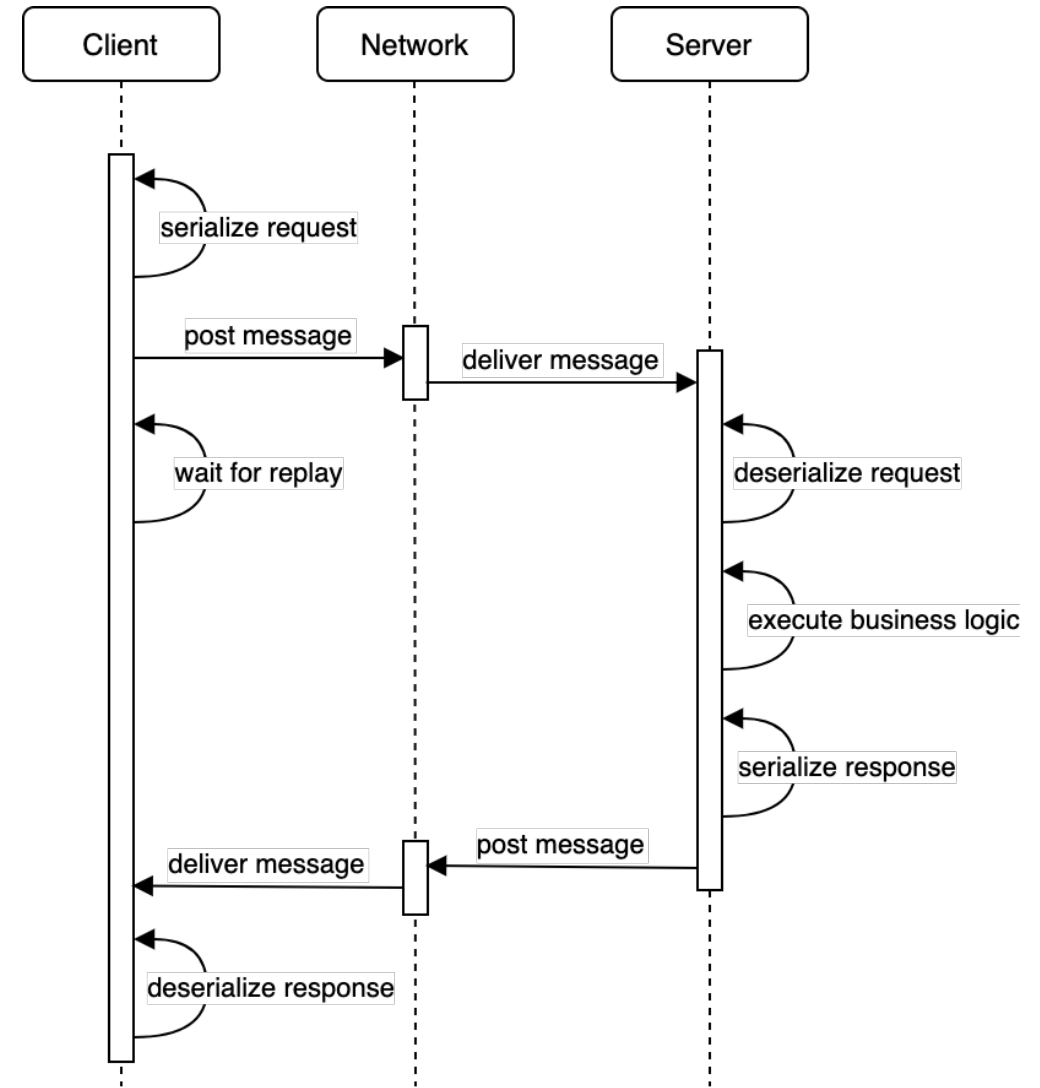
```
sayHello("Lukasz");  
  
private String sayHello(String name) {  
    return String.format(  
        "Greetings %s!", name  
    );  
}
```

Steps in RPC



Faults in RPC

- What are the consequences of?
 - Client fails to serialize request or send message to the network
 - Server fails to deserialize request e.g., invalid format of the message
 - Server raises error during execution of business logic
 - Network fails to deliver reply to the client or server takes too long to respond (timeout)




RPC Error Handling

```
(error, reply) = network.send(remote, actionData)
switch error
  case POST_FAILED:
    // handle case where you know server didn't get it
  case RETRYABLE:
    // handle case where server got it but reported transient failure
  case FATAL:
    // handle case where server got it and definitely doesn't like it
  case UNKNOWN: // i.e., time out
    // handle case where the *only* thing you know is that the server received
    // the message; it may have been trying to report SUCCESS, FATAL, or RETRYABLE
  case SUCCESS:
    if validate(reply)
      // do something with reply object
    else
      // handle case where reply is corrupt/incompatible
```

Source: <https://aws.amazon.com/builders-library/challenges-with-distributed-systems>

Message Transport and Format

- Message transport defines a protocol used to exchange data between client and a remote server. Examples:
 - TCP or UDP sockets
 - FTP file exchange
 - HTTP
 - Messaging (e.g. JMS, Kafka, MQTT, RabbitMQ)
- Message format specifies how data is being represented. Examples:
 - XML
 - JSON
 - Google Protocol Buffers
 - Apache Avro





```
<customer>
  <firstName>Luke</firstName>
  <lastName>Skywalker</lastName>
</customer>
```

```
{
  "firstName": "Luke",
  "lastName": "Skywalker"
}
```



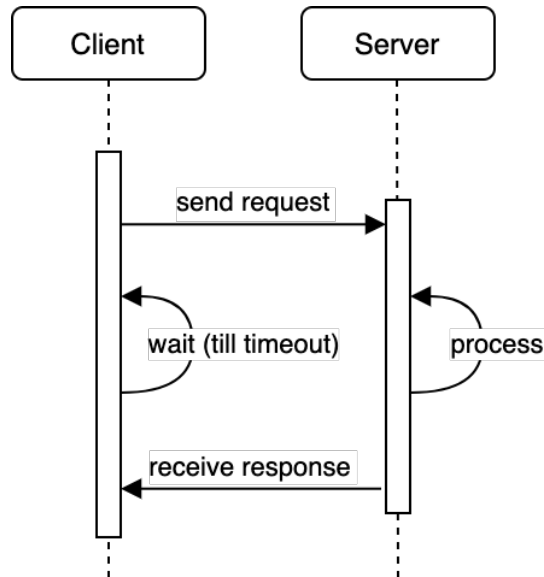
Message Transport and Format

- Performance of modern binary serializers – Avro and ProtoBuf




Source: <https://blog.softwaremill.com/the-best-serialization-strategy-for-event-sourcing-9321c299632b>

Synchronous (Request-Replay)



- Use for short running processes, without human intervention
- Always define timeout at client side
- Implement idempotent services for safe retry upon timeout

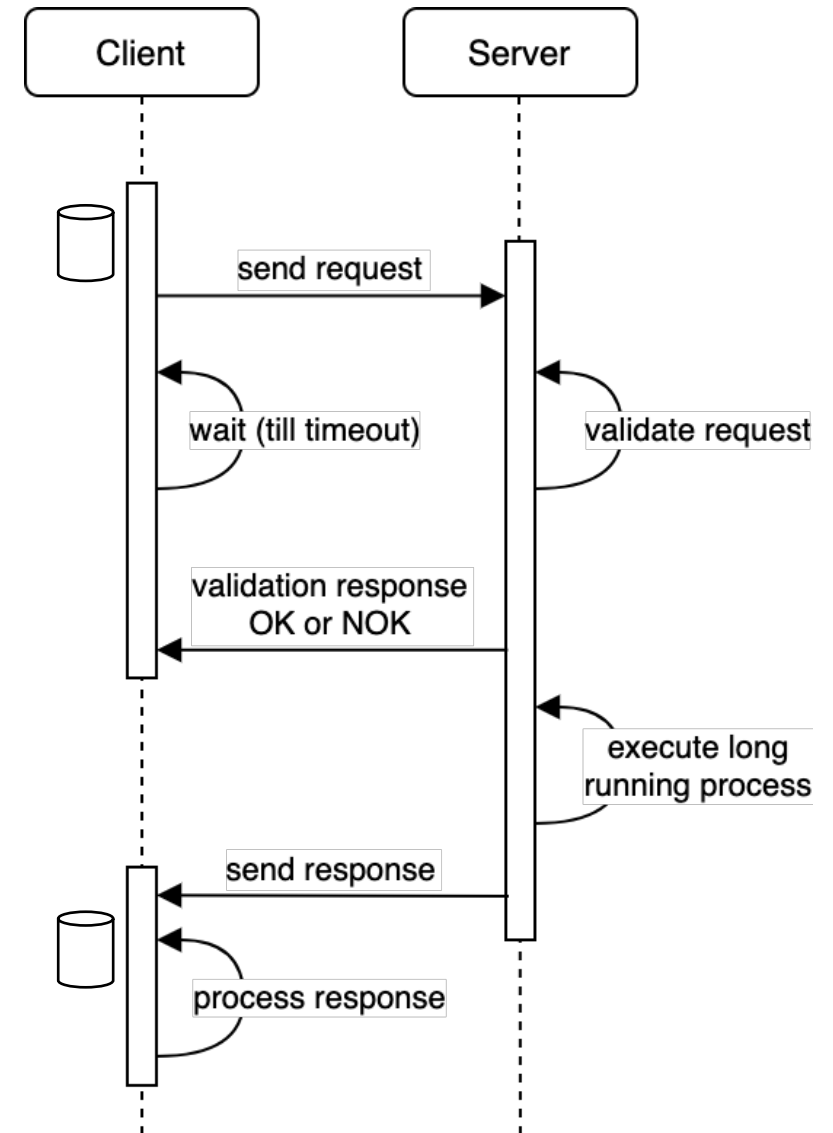
Asynchronous (Fire-and-Forget)



- Use for long running processes or when human intervention required
- Usually requires messaging solution to decouple service consumer from service provider


Communication Patterns

- Asynchronous with Confirmation
 - Server provides synchronous response only after quick validation of client's request
 - Final feedback is sent after completing long running process
- Use correlation ID in asynchronous communication
 - Client generates unique correlation ID and includes it in request message
 - Client stores request context and correlation ID in the database
 - Server provides response with given correlation ID
 - Client looks up request context by correlation ID in the database



Traditional Load Balancers


- Hardware (or software) load balancers distribute traffic between configured number of backend servers
- If the backend server is unavailable, requests are forwarded to other servers from the pool
- Backend servers should expose “health-check” endpoints that will be invoked by load balancer every N seconds to confirm their availability
 - Verify availability of dependent components in health-check implementation



Traditional Load Balancers

Apache Web Server or NGINX can be easily configured as software load balancer

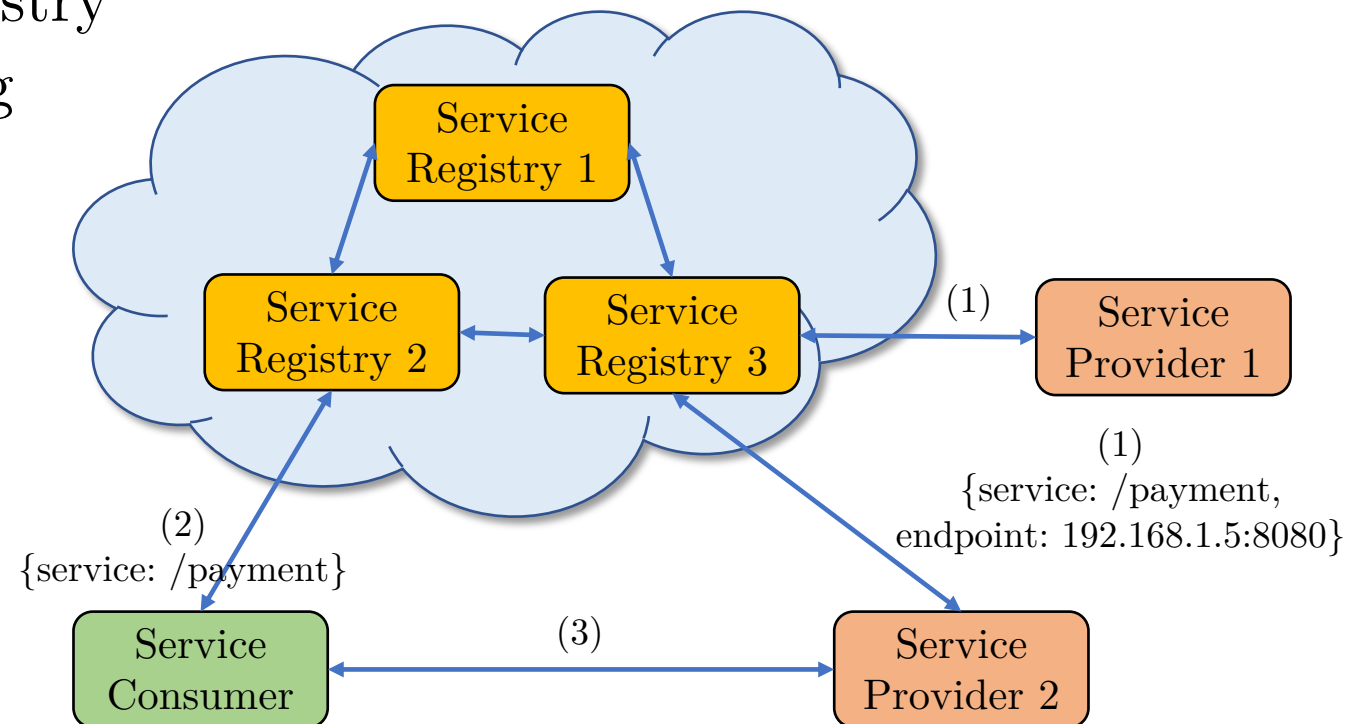
```
http {
    upstream myApp1 {
        server srv1.example.com;
        server srv2.example.com;
        server srv3.example.com;
    }
    server {
        listen 80;
        location / {
            proxy_pass http://myApp1;
            health_check interval=3 fails=3 passes=2 uri=/health;
        }
    }
}
```



Single Point of Failure!

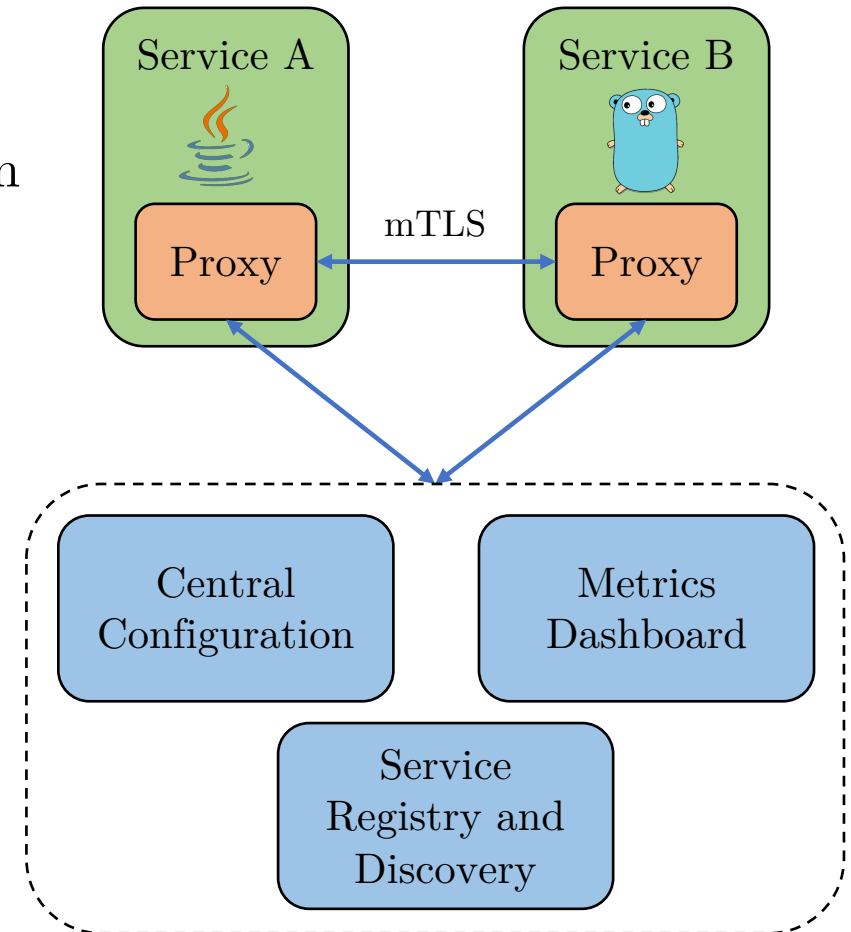
Service Registry and Discovery

- Service producers register in distributed Service Registry (1)
 - Service Registry keeps track of all alive instances of given service
 - Each instance of service provider periodically heartbeats with registry
 - Easy dynamic scaling (up and down) of service producers
- Clients query any of Service Registry nodes to get all endpoints hosting given service (2)
- Clients contact service providers directly (3)
 - Client-side load balancing based on various metrics
 - No extra hops between client and a service



Service Mesh

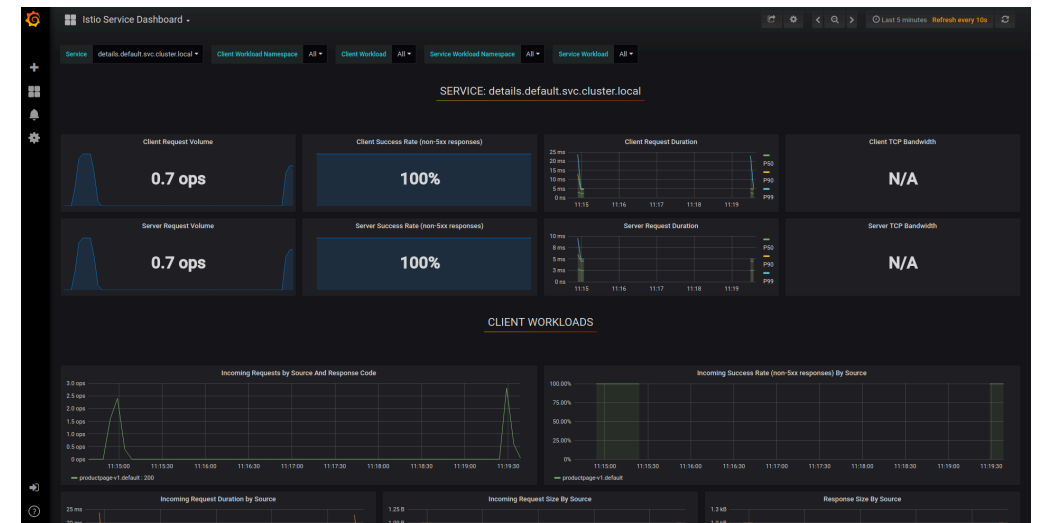
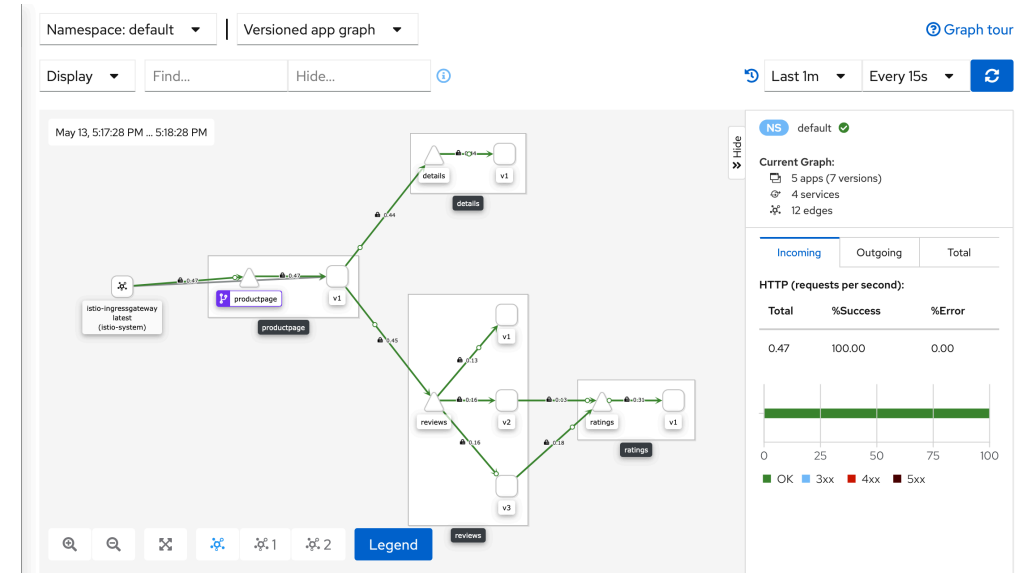
- Leverage Proxy Sidecar pattern to abstract inter-service communication
- Service Mesh features:
 - Service discovery and registry
 - Request load balancing and routing
 - Retry logic, rate limiting and Circuit Breaker pattern
 - Transport-level security (mTLS) and ACLs
 - Performance metrics
 - Distributed track and trace
 - Traffic splitting and Canary Deployments pattern
- Programming language independent inter-communication
- Example frameworks:
 - Linkerd 1.x (standalone) and 2.x (Kubernetes)
 - Istio



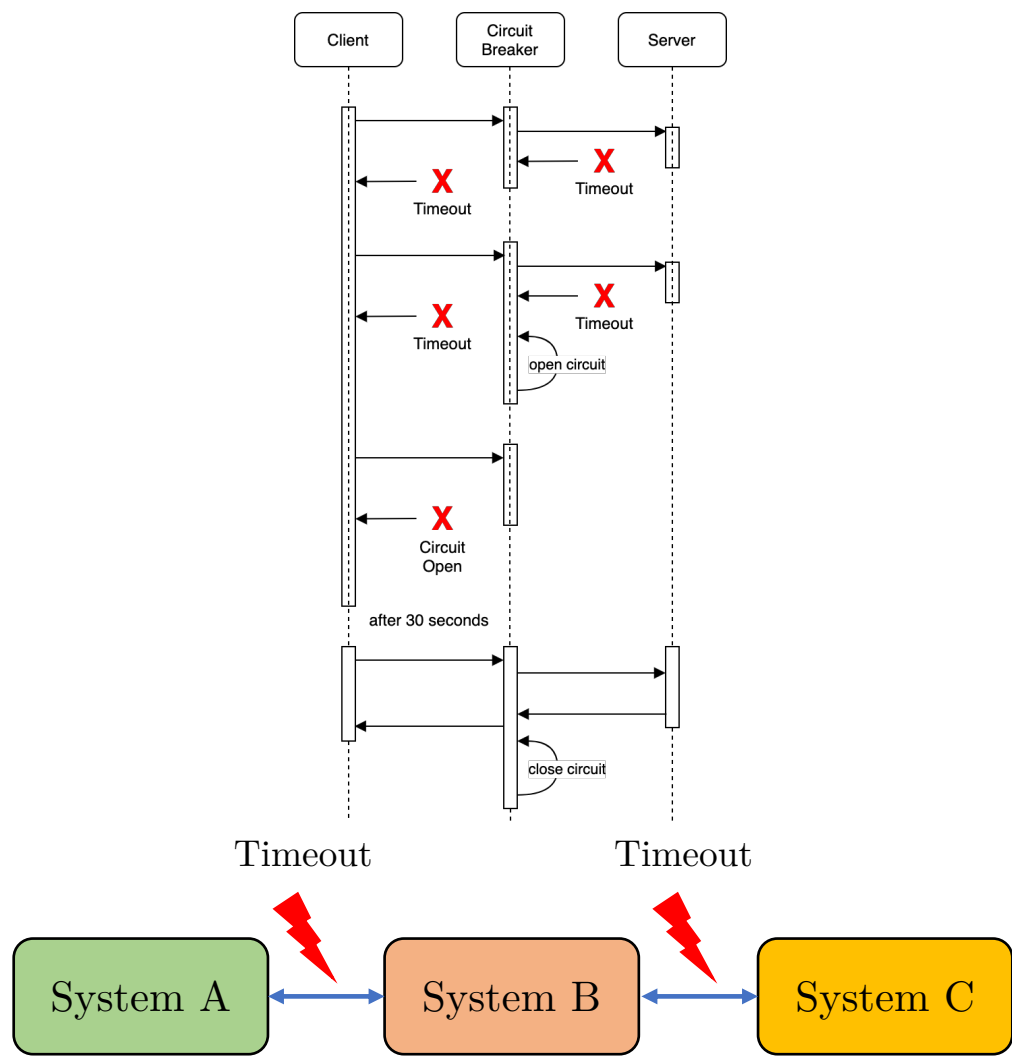
Service Mesh

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: jwt-per-host
  namespace: istio-system
spec:
  selector:
    matchLabels:
      istio: ingressgateway
  action: ALLOW
  rules:
    - from:
        - source:
            requestPrincipals: ["*@example.com"]
      to:
        - operation:
            hosts: ["example.com", "*.example.com"]
```


Source: [Istio Documentation](#)




- Circuit Breaker



- Canary Deployment



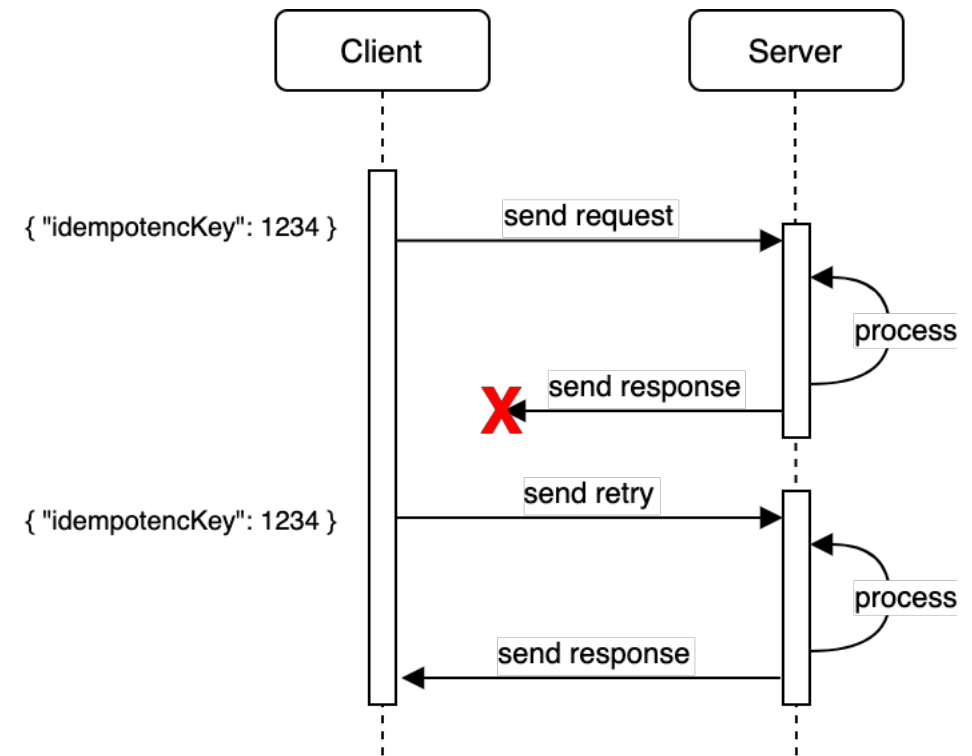
- Distributed Tracing



Definition

An idempotent service can be called many times without different outcomes, provided all requests share the same, unique idempotency key (generated by service consumer).

- Benefits:
 - Handling of duplicate requests
 - Clients have the possibility to retry timed out requests safely



Idempotent Service Design

- Verification scenarios:
 - Any operation within service implementation fails, customer retries
 - Complete service invocation succeeds, but consumer times out and retries
 - Two identical requests picked up at the same time by the system

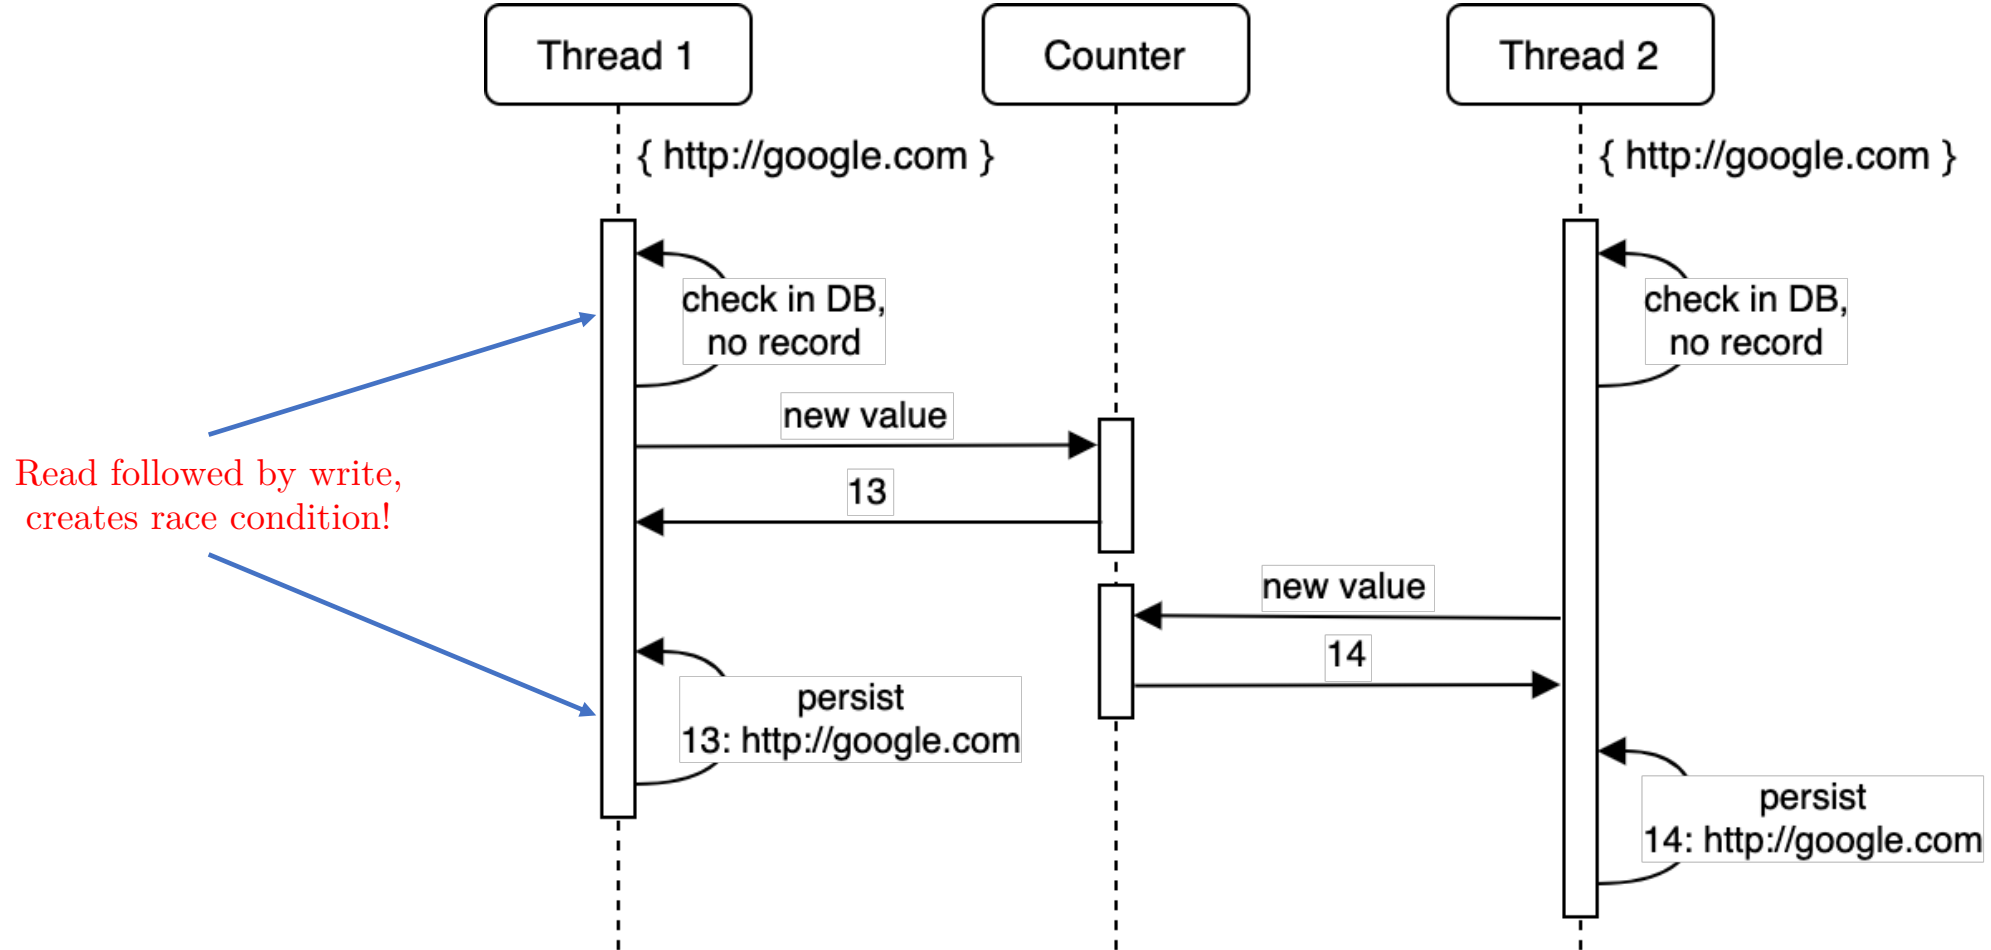
```
public String shortenUrl(String longUrl) {  
    long id = generateNextId();  
    String shortUrl = String.format(  
        "http://tiny.com:8080/%s",  
        Long.toString(id, 36));  
    insertMapping(shortUrl, longUrl);  
    return shortUrl;  
}
```

Loose one ID number

Nothing happens, local operations only
Orphaned mapping inserted
Client will retry the request

Idempotent Service Design

- Two identical requests picked up at the same time by the system



- Invoking non-idempotent services
 - Keep track whether we tried to invoke given service with idempotency key
 - We cannot retry the service invocation for the second time
 - Raise alert for operations team if we have called the service before

```
if (! didInvokeBefore("serviceA", idempotencyKey)) {
    persistStatus("serviceA", TPC_INVOKE, idempotencyKey);
    response = invokeServiceA(idempotencyKey, request);
    updateStatus("serviceA", INVOKED, idempotencyKey);
}

public boolean didInvokeBefore(String serviceName, String key) {
    Status status = queryStatus(serviceName, key);
    if (TPC_INVOKE.equals(status)) throw new UnsupportedOperationException(serviceName, key);
    return INVOKED.equals(status);
}
```

- Client and server communicate over the network
- Message transport and format
- Request-replay and Fire-and-forget communication patterns
- Service Registry and Discovery
- Service Meshes
- Idempotent service design